

## **Lecture 7**

### **Floating Point Arithmetic**

# Recap

---

- Operations on integers
  - ◆ Addition and subtraction
  - ◆ Multiplication and division

# Outline

---

- Floating point representation
- Floating point addition
- Floating point multiplication
- MIPS floating point instructions
- Subword parallelism

# Floating Point

- Representation for non-integral numbers

- ◆ Including very small and very large numbers

- Like scientific notation

- ◆  $-2.34 \times 10^{56}$

← normalized

- ◆  $+0.002 \times 10^{-4}$

← not normalized

- ◆  $+987.02 \times 10^9$

←

- In binary

- ◆  $\pm 1.xxxxxxx_2 \times 2^{yyyy}$

- Types `float` and `double` in C

# Floating Point Standard

- Defined by IEEE Std 754-1985
- Developed in response to divergence of representations
  - ◆ Portability issues for scientific code
- Now almost universally adopted
- Two representations
  - ◆ Single precision (32-bit)
  - ◆ Double precision (64-bit)

# IEEE Floating-Point Format

$$\pm 1.xxxxxxx_2 \times 2^{yyyy}$$

single: 8 bits

single: 23 bits

double: 11 bits

double: 52 bits

S	Exponent (yyyy+Bias)	Fraction (xxxx)
---	----------------------	-----------------

$$x = (-1)^S \times (1 + \text{Fraction}) \times 2^{(\text{Exponent} - \text{Bias})}$$

- S: sign bit (0  $\Rightarrow$  non-negative, 1  $\Rightarrow$  negative)
- Normalize significant:  $1.0 \leq |\text{significant}| < 2.0$ 
  - ◆ Always has a leading pre-binary-point 1 bit, so no need to represent it explicitly (hidden bit)
  - ◆ Significant is Fraction with the “1.” restored
- Exponent: excess representation: actual exponent + Bias
  - ◆ Ensures exponent is unsigned
  - ◆ Single: Bias = 127; Double: Bias = 1203

# Single-Precision Range

- Exponents 00000000 and 11111111 reserved
- Smallest value
  - ◆ Exponent: 00000001  
 $\Rightarrow$  actual exponent =  $1 - 127 = -126$
  - ◆ Fraction: 000...00  $\Rightarrow$  significand = 1.0
  - ◆  $\pm 1.0 \times 2^{-126} \approx \pm 1.2 \times 10^{-38}$
- Largest value
  - ◆ exponent: 11111110  
 $\Rightarrow$  actual exponent =  $254 - 127 = +127$
  - ◆ Fraction: 111...11  $\Rightarrow$  significand  $\approx 2.0$
  - ◆  $\pm 2.0 \times 2^{+127} \approx \pm 3.4 \times 10^{+38}$

# Double-Precision Range

- Exponents 0000...00 and 1111...11 reserved
- Smallest value
  - ◆ Exponent: 000000000001  
 $\Rightarrow$  actual exponent =  $1 - 1023 = -1022$
  - ◆ Fraction: 000...00  $\Rightarrow$  significand = 1.0
  - ◆  $\pm 1.0 \times 2^{-1022} \approx \pm 2.2 \times 10^{-308}$
- Largest value
  - ◆ Exponent: 111111111110  
 $\Rightarrow$  actual exponent =  $2046 - 1023 = +1023$
  - ◆ Fraction: 111...11  $\Rightarrow$  significand  $\approx 2.0$
  - ◆  $\pm 2.0 \times 2^{+1023} \approx \pm 1.8 \times 10^{+308}$



# Floating-Point Precision

- Relative precision

- ◆ all fraction bits are significant

- ◆  $\Delta A/|A| = 2^{-23} \times 2^{\text{exponent}} / |1 \times 2^{\text{exponent}}| = 2^{-23}$

- ◆ Single: approx  $2^{-23}$

- Equivalent to  $23 \times \log_{10} 2 \approx 23 \times 0.3 \approx 6$  decimal digits of precision

- ◆ Double: approx  $2^{-52}$

- Equivalent to  $52 \times \log_{10} 2 \approx 52 \times 0.3 \approx 16$  decimal digits of precision

# Floating-Point Example

- Represent  $-0.75$

- ◆  $-0.75 = (-1)^1 \times 1.1_2 \times 2^{-1}$

- ◆  $S = 1$

- ◆ Fraction =  $1000...00_2$

- ◆ Exponent =  $-1 + \text{Bias}$

- Single:  $-1 + 127 = 126 = 01111110_2$

- Double:  $-1 + 1023 = 1022 = 01111111110_2$

- Single:  $1011111101000...00$

- Double:  $1011111111101000...00$

# Floating-Point Example

- What number is represented by the single-precision float

11000000101000...00

- ◆  $S = 1$

- ◆ Fraction =  $01000...00_2$

- ◆ Exponent =  $10000001_2 = 129$

- $$\begin{aligned}x &= (-1)^1 \times (1 + 01_2) \times 2^{(129 - 127)} \\&= (-1) \times 1.25 \times 2^2 \\&= -5.0\end{aligned}$$

# Floating-Point Addition

- Consider a 4-digit decimal example
  - ◆  $9.999 \times 10^1 + 1.610 \times 10^{-1}$
- 1. Align decimal points
  - ◆ Shift number with smaller exponent
  - ◆  $9.999 \times 10^1 + 0.016 \times 10^1$
- 2. Add significands
  - ◆  $9.999 \times 10^1 + 0.016 \times 10^1 = 10.015 \times 10^1$
- 3. Normalize result & check for over/underflow
  - ◆  $1.0015 \times 10^2$
- 4. Round and renormalize if necessary
  - ◆  $1.002 \times 10^2$

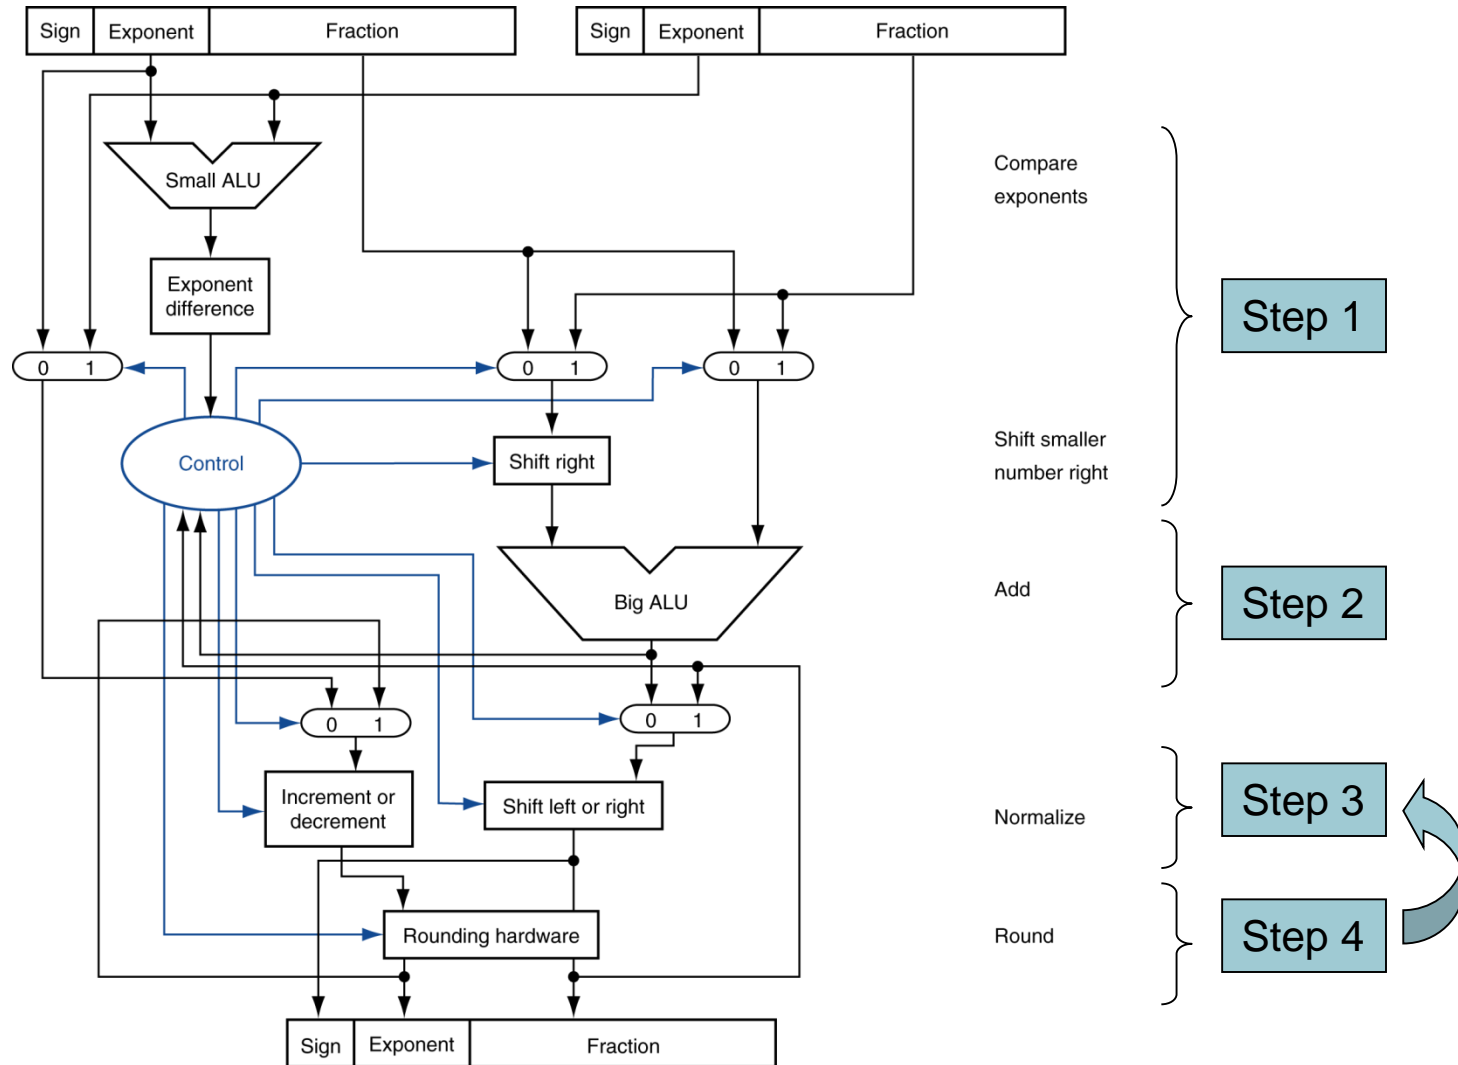
# Floating-Point Addition

- Now consider a 4-digit binary example
  - ◆  $1.000_2 \times 2^{-1} + -1.110_2 \times 2^{-2} \text{ (0.5 + -0.4375)}$
- 1. Align binary points
  - ◆ Shift number with smaller exponent
  - ◆  $1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1}$
- 2. Add significands
  - ◆  $1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1} = 0.001_2 \times 2^{-1}$
- 3. Normalize result & check for over/underflow
  - ◆  $1.000_2 \times 2^{-4}$ , with no over/underflow
- 4. Round and renormalize if necessary
  - ◆  $1.000_2 \times 2^{-4} \text{ (no change)} = 0.0625$

# FP Adder Hardware

- Much more complex than integer adder
- Doing it in one clock cycle would take too long
  - ◆ Much longer than integer operations
  - ◆ Slower clock would penalize all instructions
- FP adder usually takes several cycles
  - ◆ Can be pipelined

# FP Adder Hardware



# FP Multiplication

- Similar steps:
  - ◆ Compute exponent (careful!)
  - ◆ Multiply significands (set the binary point correctly)
  - ◆ Normalize
  - ◆ Round (potentially re-normalize)
  - ◆ Assign sign



# FP Arithmetic Hardware

- FP arithmetic hardware usually does
  - ◆ Addition, subtraction, multiplication, division, reciprocal, square-root
  - ◆  $\text{FP} \leftrightarrow \text{integer}$  conversion
- Operations usually takes several cycles
  - ◆ Can be pipelined

# FP Instructions in MIPS

- FP hardware is coprocessor 1
  - ◆ Adjunct processor that extends the ISA
- Separate FP registers
  - ◆ 32 single-precision: \$f0, \$f1, ... \$f31
  - ◆ Paired for double-precision: \$f0/\$f1, \$f2/\$f3, ...
    - Release 2 of MIPS ISA supports  $32 \times 64$ -bit FP reg's
- FP instructions operate only on FP registers
  - ◆ Programs generally don't do integer ops on FP data, or vice versa
  - ◆ More registers with minimal code-size impact
- FP load and store instructions
  - ◆ lwc1, ldc1, swc1, sdc1
    - e.g., ldc1 \$f8, 32(\$sp)

# FP Instructions in MIPS

- Single-precision arithmetic
  - ◆ `add.s`, `sub.s`, `mul.s`, `div.s`
    - e.g., `add.s $f0, $f1, $f6`
- Double-precision arithmetic
  - ◆ `add.d`, `sub.d`, `mul.d`, `div.d`
    - e.g., `mul.d $f4, $f4, $f6`
- Single- and double-precision comparison
  - ◆ `c.xx.s`, `c.xx.d` (`xx` is `eq`, `lt`, `le`, ...)
  - ◆ Sets or clears FP condition-code bit
    - e.g. `c.lt.s $f3, $f4`
- Branch on FP condition code true or false
  - ◆ `bc1t`, `bc1f`
    - e.g., `bc1t TargetLabel`

# FP Example: ° F to ° C

- C code:

```
float f2c (float fahr) {  
    return ((5.0/9.0)*(fahr - 32.0));  
}
```

- ◆ fahr in \$f12, result in \$f0, literals in global memory space

- Compiled MIPS code:

```
f2c: lwc1    $f16, const5($gp)  
     lwc1    $f18, const9($gp)  
     div.s   $f16, $f16, $f18  
     lwc1    $f18, const32($gp)  
     sub.s   $f18, $f12, $f18  
     mul.s   $f0, $f16, $f18  
     jr      $ra
```

# FP Example: Array Multiplication

- $X = X + Y \times Z$ 
  - ◆ All  $32 \times 32$  matrices, 64-bit double-precision elements

- C code:

```
void mm (double x[][],
         double y[][], double z[][]) {
    int i, j, k;
    for (i = 0; i != 32; i = i + 1)
        for (j = 0; j != 32; j = j + 1)
            for (k = 0; k != 32; k = k + 1)
                x[i][j] = x[i][j]
                    + y[i][k] * z[k][j];
}
```

- ◆ Addresses of x, y, z in \$a0, \$a1, \$a2, and  
i, j, k in \$s0, \$s1, \$s2

# FP Example: Array Multiplication

## ■ MIPS code:

	li	\$t1, 32	# \$t1 = 32 (row size/loop end)
	li	\$s0, 0	# i = 0; initialize 1st for loop
L1:	li	\$s1, 0	# j = 0; restart 2nd for loop
L2:	li	\$s2, 0	# k = 0; restart 3rd for loop
	sll	\$t2, \$s0, 5	# \$t2 = i * 32 (size of row of x)
	addu	\$t2, \$t2, \$s1	# \$t2 = i * size(row) + j
	sll	\$t2, \$t2, 3	# \$t2 = byte offset of [i][j]
	addu	\$t2, \$a0, \$t2	# \$t2 = byte address of x[i][j]
	l.d	\$f4, 0(\$t2)	# \$f4 = 8 bytes of x[i][j]
L3:	sll	\$t0, \$s2, 5	# \$t0 = k * 32 (size of row of z)
	addu	\$t0, \$t0, \$s1	# \$t0 = k * size(row) + j
	sll	\$t0, \$t0, 3	# \$t0 = byte offset of [k][j]
	addu	\$t0, \$a2, \$t0	# \$t0 = byte address of z[k][j]
	l.d	\$f16, 0(\$t0)	# \$f16 = 8 bytes of z[k][j]

...

# FP Example: Array Multiplication

...

sll	\$t0, \$s0, 5	# \$t0 = i*32 (size of row of y)
addu	\$t0, \$t0, \$s2	# \$t0 = i*size(row) + k
sll	\$t0, \$t0, 3	# \$t0 = byte offset of [i][k]
addu	\$t0, \$a1, \$t0	# \$t0 = byte address of y[i][k]
l.d	\$f18, 0(\$t0)	# \$f18 = 8 bytes of y[i][k]
mul.d	\$f16, \$f18, \$f16	# \$f16 = y[i][k] * z[k][j]
add.d	\$f4, \$f4, \$f16	# f4=x[i][j] + y[i][k]*z[k][j]
addiu	\$s2, \$s2, 1	# \$k k + 1
bne	\$s2, \$t1, L3	# if (k != 32) go to L3
s.d	\$f4, 0(\$t2)	# x[i][j] = \$f4
addiu	\$s1, \$s1, 1	# \$j = j + 1
bne	\$s1, \$t1, L2	# if (j != 32) go to L2
addiu	\$s0, \$s0, 1	# \$i = i + 1
bne	\$s0, \$t1, L1	# if (i != 32) go to L1

# Accurate Arithmetic

- IEEE Std 754 specifies additional rounding control
  - ◆ Extra bits of precision (guard, round, sticky)
  - ◆ Choice of rounding modes
  - ◆ Allows programmer to fine-tune numerical behavior of a computation
- Example:
  - ◆  $2.56 + 234$
  - ◆ 2.37 vs. 2.36
- Trade-off between hardware complexity, performance, and market requirements



# Subword Parallelism

- Graphics and audio applications can take advantage of performing simultaneous operations on short vectors
  - ◆ Example: 128-bit adder:
    - Sixteen 8-bit adds
    - Eight 16-bit adds
    - Four 32-bit adds
- Also called data-level parallelism, vector parallelism, or Single Instruction, Multiple Data (SIMD)

# Streaming SIMD Extension 2 (SSE2)

- Adds 4  $\times$  128-bit registers
  - ◆ Extended to 8 registers in AMD64/EM64T
- Can be used for multiple FP operands
  - ◆ 2  $\times$  64-bit double precision
  - ◆ 4  $\times$  32-bit double precision
  - ◆ Instructions operate on them simultaneously
    - Single-Instruction Multiple-Data

# Matrix Multiply

## ■ Unoptimized code:

```
1. void dgemm (int n, double* A, double* B, double* C)
2. {
3.     for (int i = 0; i < n; ++i)
4.         for (int j = 0; j < n; ++j)
5.             {
6.                 double cij = C[i+j*n]; /* cij = C[i][j] */
7.                 for(int k = 0; k < n; k++ )
8.                     cij += A[i+k*n] * B[k+j*n]; /* cij += A[i][k]*B[k][j] */
9.                 C[i+j*n] = cij; /* C[i][j] = cij */
10.            }
11. }
```

# Matrix Multiply

## ■ x86 assembly code:

```
1. vmovsd (%r10),%xmm0    # Load 1 element of C into %xmm0
2. mov %rsi,%rcx          # register %rcx = %rsi
3. xor %eax,%eax          # register %eax = 0
4. vmovsd (%rcx),%xmm1    # Load 1 element of B into %xmm1
5. add %r9,%rcx           # register %rcx = %rcx + %r9
6. vmulsd (%r8,%rax,8),%xmm1,%xmm1 # Multiply %xmm1,
   element of A
7. add $0x1,%rax          # register %rax = %rax + 1
8. cmp %eax,%edi          # compare %eax to %edi
9. vaddsd %xmm1,%xmm0,%xmm0 # Add %xmm1, %xmm0
10. jg 30 <dgemm+0x30>    # jump if %eax > %edi
11. add $0x1,%r11d        # register %r11 = %r11 + 1
12. vmovsd %xmm0, (%r10)  # Store %xmm0 into C element
```

# Matrix Multiply

## ■ Optimized C code:

```
1. #include <x86intrin.h>
2. void dgemm (int n, double* A, double* B, double* C)
3. {
4.     for ( int i = 0; i < n; i+=4 )
5.         for ( int j = 0; j < n; j++ ) {
6.             __m256d c0 = _mm256_load_pd(C+i+j*n); /* c0 = C[i][j]
               */
7.             for( int k = 0; k < n; k++ )
8.                 c0 = _mm256_add_pd(c0, /* c0 += A[i][k]*B[k][j] */
9.                                     _mm256_mul_pd(_mm256_load_pd(A+i+k*n),
10.                                                    _mm256_broadcast_sd(B+k+j*n)));
11.             _mm256_store_pd(C+i+j*n, c0); /* C[i][j] = c0 */
12.         }
13. }
```

# Matrix Multiply

## ■ Optimized x86 assembly code:

```
1. vmovapd (%r11),%ymm0      # Load 4 elements of C into %ymm0
2. mov %rbx,%rcx             # register %rcx = %rbx
3. xor %eax,%eax             # register %eax = 0
4. vbroadcastsd (%rax,%r8,1),%ymm1 # Make 4 copies of B element
5. add $0x8,%rax             # register %rax = %rax + 8
6. vmulpd (%rcx),%ymm1,%ymm1 # Parallel mul %ymm1, 4 A elements
7. add %r9,%rcx              # register %rcx = %rcx + %r9
8. cmp %r10,%rax             # compare %r10 to %rax
9. vaddpd %ymm1,%ymm0,%ymm0  # Parallel add %ymm1, %ymm0
10. jne 50 <dgemm+0x50>      # jump if not %r10 != %rax
11. add $0x1,%esi            # register % esi = % esi + 1
12. vmovapd %ymm0, (%r11)    # Store %ymm0 into 4 C elements
```

# Concluding Remarks

- Bits have no inherent meaning
  - ◆ Interpretation depends on the instructions applied
- Computer representations of numbers
  - ◆ Finite range and precision
  - ◆ Need to account for this in programs
- ISAs support arithmetic
  - ◆ Signed and unsigned integers
  - ◆ Floating-point approximation to reals
- Bounded range and precision
  - ◆ Operations can overflow and underflow

# Homework

---

- Due on April 9 23:55pm
- 3.27, 3.29, 3.30.